# [High Priority] Q7 Q6 Q5 Q3 Q2 [Low Priority]

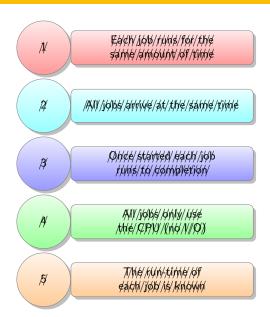
# CS 250

OPERATING SYSTEMS

Lecture 5 The Multi-Level Feedback Queue

Instructor
Dr. Dhiman Saha

# Workload Assumptions



#### Recall

- ▶ Interactive jobs care about response time
- ▶ Batch jobs care about turnaround time

## Scheduling without perfect knowledge

How can we design a scheduler that both minimizes response time for interactive jobs while also minimizing turnaround time **without** a priori knowledge of job length?

## Adaptive Scheduling

#### Recall

Algorithms like Round Robin **reduce response time** but are terrible for turnaround time. Why?

► How can the scheduler learn, as the system runs, the characteristics of the jobs it is running, and thus make better scheduling decisions?

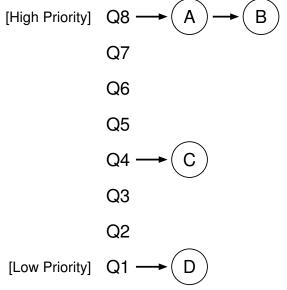
Idea

Multiple levels of round-robin

- ▶ Due to Corbato et al. [1962]
- ► Later awarded Turing Award
- ► MLFQ has a number of distinct queues
- ► Each assigned a different **priority level**
- MLFQ uses priorities to decide which job should run at a given time

A job with higher priority is chosen to run.

Rule 1 If Priority(A) > Priority(B), A runs (B doesn't).



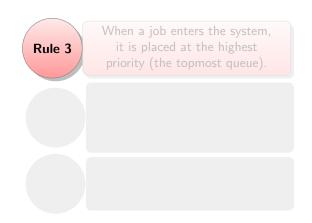
## How to set priorities?

The key to MLFQ scheduling therefore lies in how the scheduler sets priorities

- ► Approach 1: User supplied. e.g. nice
- ► Approach 2: Observed behavior (learn from history)

Rather than giving a fixed priority to each job, MLFQ varies the priority of a job based on its observed behavior

# How To Change Priority?



When a job enters the system, it is placed at the highest priority (the topmost queue).

Rule 4a

If a job uses up an entire time slice while running, its priority is reduced (i.e., it moves down one queue)

When a job enters the system, it is placed at the highest priority (the topmost queue).

Rule 4a

If a job uses up an entire time slice while running, its priority is reduced (i.e., it moves down one queue)

Rule 4b

If a job gives up the CPU before the time slice is up, it stays at the same priority leve

When a job enters the system, it is placed at the highest priority (the topmost queue).

Rule 4a

If a job uses up an entire time slice while running, its priority is reduced (i.e., it moves down one queue)

Rule 4b

If a job gives up the CPU before the time slice is up, it stays at the same priority leve

When a job enters the system, it is placed at the highest priority (the topmost queue).

Rule 4a

If a job uses up an entire time slice while running, its priority is reduced (i.e., it moves down one queue)

Rule 4b

If a job gives up the CPU before the time slice is up, it stays at the same priority leve

When a job enters the system, it is placed at the highest priority (the topmost queue).

Rule 4a

If a job uses up an entire time slice while running, its priority is reduced (i.e., it moves down one queue)

Rule 4b

If a job gives up the CPU before the time slice is up, it stays at the same priority level

When a job enters the system, it is placed at the highest priority (the topmost queue).

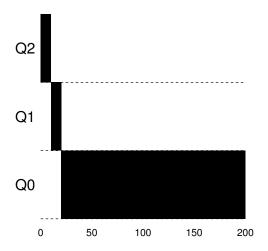
Rule 4a

If a job uses up an entire time slice while running, its priority is reduced (i.e., it moves down one queue)

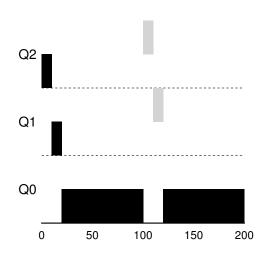
Rule 4b

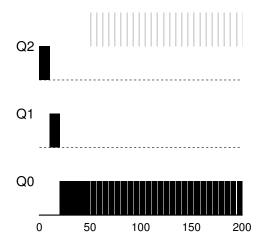
If a job gives up the CPU before the time slice is up, it stays at the same priority level

# Long-running Job Over Time



# Along Came An Interactive Job





A Mixed I/O-intensive and CPU-intensive Workload

Starvation

What if there are too many interactive jobs in the system?

#### Gaming the scheduler

What if an user tricks the scheduler into giving you more than your fair share of the CPU?

Change in behavior of a job

Is the current approach adaptive?

After some time period S, move all the jobs in the system to the topmost queue

#### Solved

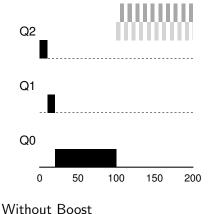
- ► Starvation
- ► Adapting to change in behavior

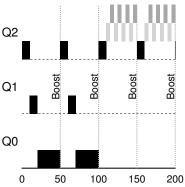
## Why?

Still Unsolved

► Gaming the scheduler.

▶ Value of S = 50ms





With Boost

## How to decide the value of S?

## The Voo-Doo Constant (coined by John Ousterhout)

- ► Too high, and long-running jobs could starve
- ► Too low, and interactive jobs may not get a proper share of the CPU.

Rule/4/a

If a job uses up an entire time slice while running, its priority is reduced (i.e., it moves down one queue)

R/Ne/46

If a job gives up the CPU before the time slice is up, it stays at the same priority level

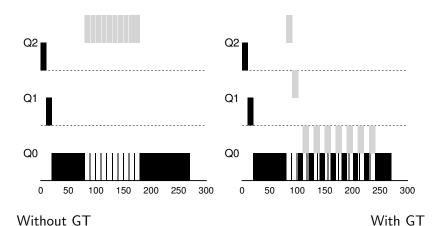
Rule 4

Once a job uses up its time allotment at a given level, its priority is reduced

## Solves Gaming Problem

- Perform better accounting of CPU time at each level of the MLFQ
- ► Keep track of actual time spent on the CPU
- ▶ Demote once allotted time is over
- Regardless of how many times the job has given up the CPU

# Gaming Tolerance (GT)



## Parameterizing MLFQ

- ► How many queues should there be?
- ► How big should the time slice be per queue?
- ► How often should priority be boosted in order to avoid starvation and account for changes in behavior?

No easy answers to these questions

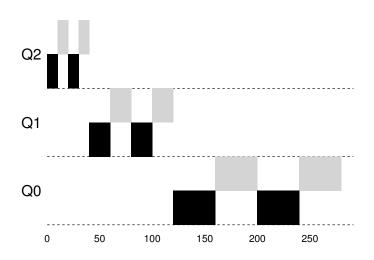
# Varying Time-slice Length

#### Varying time-slice length across different queues

- ► The high-priority queues are usually given short time slices; they are comprised of interactive jobs,
- ► The low-priority queues, in contrast, contain long-running jobs that are CPU-bound; hence, longer time slices work well

#### Other features

advice on setting priorities e.g. nice command on UNIX



If Priority(A) > Priority(B), A runs (B doesn't).

Rule 2

If Priority(A) = Priority(B), A & B run in RR

Rule 3

When a job enters the system, it is placed at the highest priority

Rule 4

Once a job uses up its time allotment at a given level, its priority is reduced

Rule 5

After some time period *S*, move all the jobs in the system to the topmost queue

Idea

#### Lottery Scheduling

#### Approach:

- Processes are alloted lottery tickets
- Whoever wins runs
- ► Higher priority implies more tickets