


CS 250

OPERATING SYSTEMS

Lecture 9 TLB

Instructor
Dr. Dhiman Saha



Space overheads

Issue 1

Storing PT in memory wastes valuable memory space.

> Factor 2 slow down

Issue 2

For every memory reference, paging requires us to perform one extra memory reference in order to first fetch the translation from the page table

Space overheads

Issue 1

Storing PT in memory wastes valuable memory space.

> Factor 2 slow down

Issue 2

For every memory reference, paging requires us to perform one extra memory reference in order to first fetch the translation from the page table

Attempt #6: Translation-lookaside Buffer

Faster Paging

- ▶ How can we speed up address translation?
- ▶ Avoid the **extra memory reference** that paging seems to require
- ▶ What hardware support is required?
- ▶ What OS involvement is needed?

- ▶ Hardware support
- ▶ Part of the chips memory-management unit
- ▶ Basically, a **hardware cache** of popular virtual-to-physical address translations
- ▶ A.k.a address-translation cache

Idea


The hardware first checks the TLB to see if the desired translation is held therein.

- ▶ Hardware support
- ▶ Part of the chips memory-management unit
- ▶ Basically, a **hardware cache** of popular virtual-to-physical address translations
- ▶ A.k.a address-translation cache

Idea

The hardware first checks the TLB to see if the desired translation is held therein.

Paging hardware with TLB




```
1  VPN = (VirtualAddress & VPN_MASK) >> SHIFT
2  (Success, TlbEntry) = TLB_Lookup(VPN)
3  if (Success == True)    // TLB Hit
4      if (CanAccess(TlbEntry.ProtectBits) == True)
5          Offset    = VirtualAddress & OFFSET_MASK
6          PhysAddr  = (TlbEntry.PFN << SHIFT) | Offset
7          Register  = AccessMemory(PhysAddr)
8      else
9          RaiseException(PROTECTION_FAULT)
10 else    // TLB Miss
11     PTEAddr = PTBR + (VPN * sizeof(PTE))
12     PTE = AccessMemory(PTEAddr)
13     if (PTE.Valid == False)
14         RaiseException(SEGMENTATION_FAULT)
15     else if (CanAccess(PTE.ProtectBits) == False)
16         RaiseException(PROTECTION_FAULT)
17     else
18         TLB_Insert(VPN, PTE.PFN, PTE.ProtectBits)
19         RetryInstruction()
```

Example

Integer Array of 10 elements

	Offset				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 02					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06		a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

What is the TLB activity pattern?

Example

Integer Array of 10 elements

	Offset				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 02					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06				a[0]	
VPN = 07	a[1]	a[2]	a[3]	a[4]	
VPN = 08	a[5]	a[6]	a[7]	a[8]	
VPN = 09	a[9]				
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

What is the TLB activity pattern?

► TLB Hit Vs TLB Miss

- Hit rate
- Effect of page-size
- Spatial Locality
- Temporal Locality

Effective Memory Access Time

How to calculate?

- ▶ TLB Hit Vs TLB Miss
- ▶ Hit rate
- ▶ Effect of page-size
- ▶ Spatial Locality
- ▶ Temporal Locality

Effective Memory Access Time

How to calculate?

- ▶ TLB Hit Vs TLB Miss
- ▶ Hit rate
- ▶ Effect of page-size
- ▶ Spatial Locality
- ▶ Temporal Locality

Effective Memory Access Time

How to calculate?

- ▶ TLB Hit Vs TLB Miss
- ▶ Hit rate
- ▶ Effect of page-size
- ▶ Spatial Locality
- ▶ Temporal Locality

Effective Memory Access Time

How to calculate?

- ▶ TLB Hit Vs TLB Miss
- ▶ Hit rate
- ▶ Effect of page-size
- ▶ Spatial Locality
- ▶ Temporal Locality

Effective Memory Access Time

How to calculate?

- ▶ TLB Hit Vs TLB Miss
- ▶ Hit rate
- ▶ Effect of page-size
- ▶ Spatial Locality
- ▶ Temporal Locality

Effective Memory Access Time

How to calculate?

Who Handles The TLB Miss?

► Hardware Vs Software

CISC

Complex-instruction set computers

- Early systems
- TLB miss handled by hardware
- via a page-table base register

RISC

Reduced-instruction set computers

- Modern architectures
- Software-managed TLB
- Hardware simply raises an exception
- Trap handling mechanism takes over

```
1  VPN = (VirtualAddress & VPN_MASK) >> SHIFT
2  (Success, TlbEntry) = TLB_Lookup(VPN)
3  if (Success == True)    // TLB Hit
4      if (CanAccess(TlbEntry.ProtectBits) == True)
5          Offset    = VirtualAddress & OFFSET_MASK
6          PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
7          Register = AccessMemory(PhysAddr)
8      else
9          RaiseException(PROTECTION_FAULT)
10 else    // TLB Miss
11     RaiseException(TLB_MISS)
```

Does the return-from-trap instruction in a TLB miss needs to be a little different than the return-from-trap when servicing a system call?

Can an infinite chain of TLB misses occur with a OS handled TLB miss?

► Possible Solutions

What are the advantages of a software-managed TLB-miss

Does the return-from-trap instruction in a TLB miss needs to be a little different than the return-from-trap when servicing a system call?

Can an infinite chain of TLB misses occur with a OS handled TLB miss?

► Possible Solutions

What are the advantages of a software-managed TLB-miss

Does the return-from-trap instruction in a TLB miss needs to be a little different than the return-from-trap when servicing a system call?

Can an infinite chain of TLB misses occur with a OS handled TLB miss?

► Possible Solutions

What are the advantages of a software-managed TLB-miss

- ▶ Typically 32, 64, 128 entries
- ▶ Fully Associative. What does this mean?

TLB Entry

VPN | PFN | other bits

Other bits

- ▶ Valid. Is this same as page valid bit (Refer next slide)?
- ▶ Protection
- ▶ ASID (Described next)
- ▶ Dirty

- ▶ Typically 32, 64, 128 entries
- ▶ Fully Associative. What does this mean?

TLB Entry

VPN | PFN | other bits

Other bits

- ▶ Valid. Is this same as page valid bit (Refer next slide)?
- ▶ Protection
- ▶ ASID (Described next)
- ▶ Dirty




- ▶ Typically 32, 64, 128 entries
- ▶ Fully Associative. What does this mean?

TLB Entry

VPN | PFN | other bits

Other bits

- ▶ Valid. Is this same as page valid bit (Refer next slide)?
- ▶ Protection
- ▶ ASID (Described next)
- ▶ Dirty



How to handle this?

- ▶ When context-switching between processes, the translations in the TLB for the last process are **not meaningful** to the about-to-be-run process.

VPN	PFN	valid	prot
10	100	1	rwX
—	—	0	—
10	170	1	rwX
—	—	0	—

How?

- ▶ Set all valid bits to 0

When

- ▶ In SW use explicit instruction
- ▶ In HW while page-table base reg is updated

Issue

- ▶ Performance overhead
- ▶ Each time a process runs, it **must incur TLB misses** as it touches its data and code pages.

How?

- ▶ Set all valid bits to 0

When

- ▶ In SW use explicit instruction
- ▶ In HW while page-table base reg is updated

Issue

- ▶ Performance overhead
- ▶ Each time a process runs, it **must incur TLB misses** as it touches its data and code pages.

How?

- ▶ Set all valid bits to 0

When

- ▶ In SW use explicit instruction
- ▶ In HW while page-table base reg is updated

Issue

- ▶ Performance overhead
- ▶ Each time a process runs, it **must incur TLB misses** as it touches its data and code pages.

How?

- ▶ Set all valid bits to 0

When

- ▶ In SW use explicit instruction
- ▶ In HW while page-table base reg is updated

Issue

- ▶ Performance overhead
- ▶ Each time a process runs, it **must incur TLB misses** as it touches its data and code pages.

VPN	PFN	valid	prot	ASID
10	100	1	rwX	1
—	—	0	—	—
10	170	1	rwX	2
—	—	0	—	—

- ▶ Analogous to PID
- ▶ Smaller in size
- ▶ With ASID the TLB can hold translations from different processes at the same time without any ambiguity.

OS-HW Support

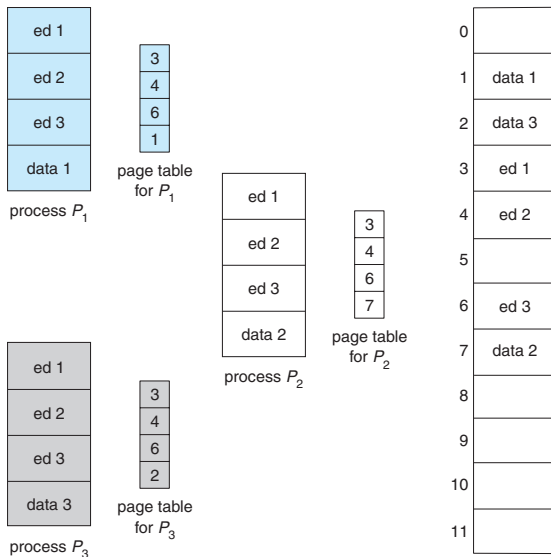
OS must, on a context switch, set some privileged register to the ASID of the current process.

VPN	PFN	valid	prot	ASID
10	100	1	rwX	1
—	—	0	—	—
10	170	1	rwX	2
—	—	0	—	—

- ▶ Analogous to PID
- ▶ Smaller in size
- ▶ With ASID the TLB can hold translations from different processes at the same time without any ambiguity.

OS-HW Support

OS must, on a context switch, set some privileged register to the ASID of the current process.



What would the same reflect in the TLB?

What would the same reflect in the TLB?

VPN	PFN	valid	prot	ASID
10	101	1	r-x	1
—	—	0	—	—
50	101	1	r-x	2
—	—	0	—	—

Which TLB entry should be replaced when we add a new TLB entry?

Goal

Minimize the miss rate (or increase hit rate)

Idea

It is likely that an entry that has **not recently been used** is a good candidate for eviction

- ▶ Takes advantage of **locality** in the memory-reference stream

Classwork

- ▶ Can you devise a program that would perform terribly with LRU?
- ▶ What would be the alternative policy?

- ▶ Evicts a TLB mapping at random
- ▶ Simple to implement
- ▶ Ability to avoid corner-case behaviors like LRU

HW-4/5

Study MIPS TLB Entry

Solves one issue with Paging

- ▶ Effective access time is decreased

However

- ▶ Issue of TLB Coverage

Does not solve

- ▶ Issue of storing large PT in memory

Solves one issue with Paging

- ▶ Effective access time is decreased

However

- ▶ Issue of TLB Coverage

Does not solve

- ▶ Issue of storing large PT in memory

Solves one issue with Paging

- ▶ Effective access time is decreased

However

- ▶ Issue of TLB Coverage

Does not solve

- ▶ Issue of storing large PT in memory

Attempt #7: Smaller Page Tables

Next Lecture